

Sample Document for RAG Testing

Retrieval Augmented Generation (RAG) is a technique used in modern AI systems to improve the accuracy of large language models. Instead of relying only on the model's training data, RAG retrieves relevant information from external documents and provides it as context to the model.

The typical RAG pipeline includes: document ingestion, text chunking, embedding generation, vector database storage, retrieval of relevant chunks, and finally answer generation using an LLM.